

# Introduction to Numerical Optimization

Janoś Gabler and Tim Mensinger, University of Bonn



estimagic

# Installation

We assume you have done the following

- Installed miniconda or anaconda
- Executed:

```
git clone https://github.com/OpenSourceEconomics/euroscipy-estimagic.git
cd euroscipy-estimagic
conda env create -f environment.yml
conda activate euroscipy-estimagic
```

- If you haven't done so, please do so until the first practice session
- Details: <https://github.com/OpenSourceEconomics/euroscipy-estimagic>

# About Us



- Website: [janosg.com](https://janosg.com)
- GitHub: [janosg](https://github.com/janosg)
- Started estimagic in 2019
- Postdoc in Econ, University of Bonn
- Open for interesting jobs



- Website: [tmensing.com](https://tmensing.com)
- GitHub: [timmens](https://github.com/timmens)
- estimagic core contributor
- PhD student in Econ, University of Bonn

# Sections

1. Introduction to **scipy.optimize**
2. Introduction to **estimagic**
3. Choosing algorithms
4. Bounds and constraints
5. Global optimization

# Structure of each topic

1. Summary of exercise you will solve
2. Some theory
3. Syntax in very simplified example
4. You solve a more difficult example in a notebook
5. Discuss one possible solution


# Introduction to `scipy.optimize`

# Preview of practice session

- Translate a criterion function from math to code
- Use **scipy.optimize** to minimize the criterion function

# Example problem

- **Criterion**  $f(a, b) = a^2 + b^2$
- Parameters  $a, b$
- Want:  $a^*, b^* = \operatorname{argmin} f(a, b)$
- Possible extensions:
  - Constraints
  - Bounds
- Optimum at  $a^* = 0, b^* = 0$ ,  
 $f(a^*, b^*) = 0$





# Optimization with `scipy.optimize`

```
>>> import numpy as np
>>> from scipy.optimize import minimize

>>> def sphere(x):
...     a, b = x
...     return a ** 2 + b ** 2

>>> x0 = np.ones(2)
>>> res = minimize(sphere, x0)
>>> res.fun
0.0
>>> res.x
array([0.0, 0.0])
```

# Features of `scipy.optimize`

- **minimize** as unified interface to 14 local optimizers
  - some support bounds
  - some support constraints
- Parameters are 1d arrays
- Maximize by minimizing  $-f(x)$
- Different interfaces for:
  - global optimization
  - nonlinear least-squares

# Practice Session

First optimization with `scipy.optimize` (15 min)

## Pros

- Very mature and reliable
- No additional dependencies
- Low overhead
- Enough algorithms for many use-cases

## Cons

- Relatively few algorithms
- No parallelization
- Maximization via sign flipping
- Feedback only at end
- No feedback in case of crash
- Parameters are flat arrays

# Examples from real projects I

```
>>> scipy.optimize.minimize(func, x0)
-----
LinAlgError                                Traceback (most recent call last)
<ipython-input-17-7459e5b4d8d4> in <module>
--> 1 scipy.optimize.minimize(func, x0)

    95
    96 def _raise_linalgerror_singular(err, flag):
--> 97     raise LinAlgError("Singular matrix")
    98

LinAlgError: Singular matrix
```

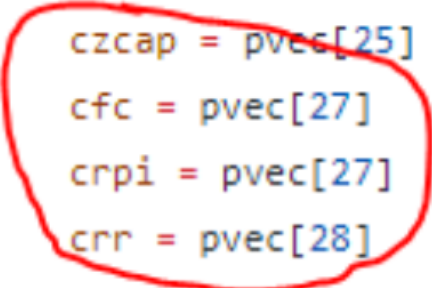
- After 5 hours and with no additional information

# Examples from real projects II

```
def parse_parameters(x):  
    """Parse the parameter vector into quantities we need."""  
    num_types = int(len(x[54:]) / 6) + 1  
    params = {  
        'delta': x[0:1],  
        'level': x[1:2],  
        'coeffs_common': x[2:4],  
        'coeffs_a': x[4:19],  
        'coeffs_b': x[19:34],  
        'coeffs_edu': x[34:41],  
        'coeffs_home': x[41:44],  
        'type_shares': x[44:44 + (num_types - 1) * 2],  
        'type_shifts': x[44 + (num_types - 1) * 2:]  
    }  
    return params
```

## Examples from real projects III

```
26      cindw = pvec[23]
27      cindp = pvec[24]
28      czcap = pvec[25]
29      cfc = pvec[27]
30      crpi = pvec[27]
31      crr = pvec[28]
32      cry = pvec[29]
```

A red hand-drawn circle highlights the code lines for cfc, crpi, and crr. The circle is drawn around the lines 29, 30, and 31, which are: cfc = pvec[27], crpi = pvec[27], and crr = pvec[28]. The circle is drawn with a red line and is slightly irregular in shape.



**Joshua Brault**  
@JoshuaBrault3



Also, in the model solution (modelsol.R) the authors parameterize the fixed cost share and the Taylor rule inflation feedback parameter to the same thing. This is obviously a typo as cfc should be parameter 26 in the vector. I have no idea how this impacts their results.

```
dajmcdon initial commit Latest commit 28d7874 on Mar 29, 2021 History
R 1 contributor

754 lines (653 sloc) 22.6 KB
Raw Blame

1 gensolution <- function(pvec) {
2   ## Function takes in a lengthy vector of parameters and returns matrices nearly ready for Kalman input
3   stderr.es = pvec[1]
4   stderr.eb = pvec[2]
5   stderr.eg = pvec[3]
6   stderr.eqs = pvec[4]
7   stderr.em = pvec[5]
8   stderr.epinf = pvec[6]
9   stderr.eu = pvec[7]
10  ##
11  crhoa = pvec[8]
12  crhob = pvec[9]
13  crhog = pvec[10]
14  crhoqs = pvec[11]
15  crhoms = pvec[12]
16  crhoqinf = pvec[13]
17  crhou = pvec[14]
18  cmap = pvec[15]
19  cmau = pvec[16]
20  csadjcost = pvec[17]
21  csigma = pvec[18]
22  chaob = pvec[19]
23  corobu = pvec[20]
24  csigl = pvec[21]
25  corobp = pvec[22]
26  cindu = pvec[23]
27  cindp = pvec[24]
28  cccap = pvec[25]
29  cfc = pvec[27]
30  crpi = pvec[27]
31  crr = pvec[28]
32  cry = pvec[29]
```



# Introduction to estimagic

# Preview of practice session

- Translate a **scipy** optimization to **estimagic.minimize**
- Use dictionaries instead of flat arrays as parameters in the optimization
- Plot the optimization history (criterion and parameter)

# What is estimagic?



estimagic

- Library for numerical optimization
- Tools for nonlinear estimation and inference
- Harmonized interface to:
  - Scipy, Nlopt, TAO, Pygmo, ...
- Adds functionality and convenience

# You can use it like scipy

```
>>> import estimagic as em
>>> def sphere(x):
...     a, b = x
...     return a ** 2 + b ** 2

>>> res = em.minimize(
...     criterion=sphere,
...     params=np.ones(2),
...     algorithm="scipy_lbfgsb",
... )

>>> res.params
array([ 0.,  0])
```

- No default algorithm
- Options have different names

# Params can be (almost) anything

```
>>> params = {"a": 0, "b": 1, "c": pd.Series([2, 3, 4])}
>>> def dict_sphere(x):
...     out = (
...         x["a"] ** 2 + x["b"] ** 2 + (x["c"] ** 2).sum()
...     )
...     return out

>>> res = em.minimize(
...     criterion=dict_sphere,
...     params=params,
...     algorithm="scipy_neldermead",
... )

>>> res.params
{'a': 0.,
 'b': 0.,
 'c': 0    0.
      1    0.
      2    0.
 dtype: float64}
```

- numpy arrays
- pd.Series, pd.DataFrame
- scalars
- (Nested) lists, dicts and tuples thereof
- Special case: DataFrame with columns **"value"**, **"lower\_bound"** and **"upper\_bound"**

# OptimizeResult

```
>>> res
```

Minimize with 5 free parameters terminated successfully after 805 criterion evaluations and 507 iterations.

The value of criterion improved from 30.0 to 1.6760003634613059e-16.

The `scipy_neldermead` algorithm reported: Optimization terminated successfully.

Independent of the convergence criteria used by `scipy_neldermead`, the strength of convergence can be assessed by the following criteria:

	one_step	five_steps
relative_criterion_change	1.968e-15***	2.746e-15***
relative_params_change	9.834e-08*	1.525e-07*
absolute_criterion_change	1.968e-16***	2.746e-16***
absolute_params_change	9.834e-09**	1.525e-08*

(\*\*\*: change  $\leq 1e-10$ , \*\*: change  $\leq 1e-8$ , \*: change  $\leq 1e-5$ . Change refers to a change between accepted steps. The first column only considers the last step. The second column considers the last five steps.)

# Access OptimizeResult's attributes

```
>>> res.criterion
.0
>>> res.n_criterion_evaluations
805
>>> res.success
True
>>> res.message
'Optimization terminated successfully.'
>>> res.history.keys():
dict_keys(['params', 'criterion', 'runtime'])
```

# Logging and Dashboard


```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="scipy_lbfgsb",  
...     logging="my_log.db",  
... )  
  
>>> from estimagic import OptimizeLogReader  
  
>>> reader = OptimizeLogReader("my_log.db")  
>>> reader.read_history().keys()  
dict_keys(['params', 'criterion', 'runtime'])  
  
>>> reader.read_iteration(1)["params"]  
array([0., 0.817, 1.635, 2.452, 3.27 ])
```

- Persistent log in sqlite database
- No data loss ever
- Can be read during optimization
- Provides data for dashboard
- No SQL knowledge needed



# Criterion plot


```
em.criterion_plot(res)
```



- First argument can be:
  - `OptimizeResult`
  - path to log file
  - list or dict thereof
- Dictionary keys are used for legend

# Criterion plot (2)


```
em.criterion_plot(res, monotone=True)
```



- **monotone=True** shows the current best value
- useful if there are extreme values in history

# Criterion plot (3)

```
em.criterion_plot(res, max_evaluations=300)
```




- **max\_evaluations** sets upper limit of x-axis

# Params plot

```
# reminder: params looks like this
params = {
    "a": 0,
    "b": 1,
    "c": pd.Series([2, 3, 4])
}

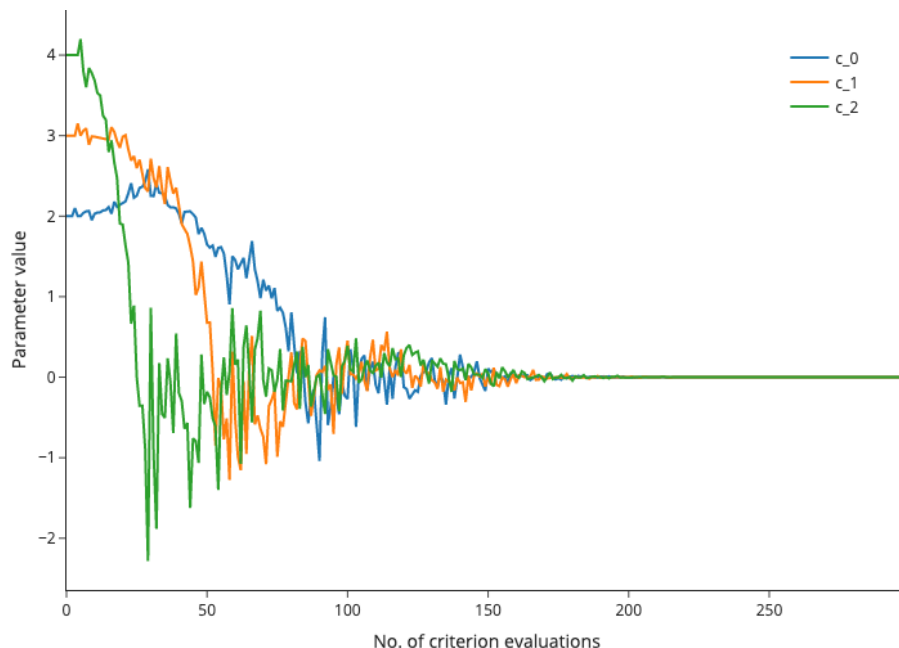
em.params_plot(
    res,
    max_evaluations=300,
)
```



- Similar options as **criterion\_plot**

# Params plot (2)

```
em.params_plot(  
    res,  
    max_evaluations=300,  
    selector=lambda x: x["c"],  
)
```



- **selector** is a function returning a subset of params

# There is maximize

```
>>> def upside_down_sphere(params):  
...     return -params @ params  
  
>>> res = em.maximize(  
...     criterion=upside_down_sphere,  
...     params=np.arange(5),  
...     algorithm="scipy_lbfgs",  
... )  
>>> res.params  
array([ 0.,  0.,  0.,  0.,  0.]
```

# Harmonized algo\_options

```
>>> algo_options = {  
...     "convergence.relative_criterion_tolerance": 1e-9,  
...     "stopping.max_iterations": 100_000,  
...     "trustregion.initial_radius": 10.0,  
... }  
  
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="nag_pybobyqa",  
...     algo_options=algo_options,  
... )  
>>> res.params  
array([0., 0., 0., 0., 0.]
```

- The same options have the same name
- Different options have different names
  - e.g., not one **tol**
- Ignore irrelevant options

# estimagic.readthedocs.io

estimagic

[Getting Started](#) [How-to Guides](#) [Explanations](#) [API](#) [Development](#) [Optimizers](#)

Previous topic

[estimagic](#)

Next topic


[Installation](#)

Quick search

Go


## Getting Started

This section contains quickstart guides for new estimagic users. It can also serve as a reference for more experienced users.



### Optimization

Learn numerical optimization with  
estimagic



### Differentiation

Learn numerical differentiation with  
estimagic



### Estimation

Learn maximum likelihood and  
methods of simulated moments  
estimation with estimagic



### Installation

Installation instructions for estimagic  
and optional dependencies



# Who is behind estimagic



Janoš Gabler



Mariam Petrosyan



Tim Mensinger



Klara Röhl

NUMFOCUS

OPEN CODE = BETTER SCIENCE




Tobias Raabe



Annica Gehlen



Sebastian Gsell



Bahar Coskun



Aida  
Takhmazova



Hans-Martin von  
Gaudecker



Kenneth L. Judd



Break (5 min)

# Practice Session

Convert previous example to estimagic (15 min)

# Choosing algorithms

# Preview of practice session

You will get an optimization problem that fails

- Figure out why it fails
- Choose an algorithm that works

# Relevant problem properties

- **Smoothness:** Differentiable? Kinks? Discontinuities? Stochastic?
  - **Convexity:** Are there local optima?
  - **Goal:** Do you need a global solution? How precise?
  - **Size:** 2 parameters? 10? 100? 1000? More?
  - **Constraints:** Bounds? Linear constraints? Nonlinear constraints?
  - **Structure:** Nonlinear least-squares, Log-likelihood function
- > Properties guide selection but experimentation is important

# scipy\_lbfgsb

- Limited memory BFGS
- BFGS: Approximate hessians from multiple gradients
- Criterion must be differentiable
- Scales to a few thousand parameters
- Beats other BFGS implementations in many benchmarks
- Low overhead

# fides

- Derivative based trust-region algorithm
- Developed by Fabian Fröhlich as a Python package
- Many advanced options to customize the optimization!
- Criterion must be differentiable
- Good solution if **scipy\_lbfgsb** is too aggressive



# nlopt\_bobyqa, nag\_pybobyqa

- **B**ound **O**ptimization **by Q**uadratic **A**pproximation
- Derivative free trust region algorithm
- **nlopt** has less overhead
- **nag** has options to deal with noise
- Good for non-differentiable not too noisy functions
- Slower than derivative based methods but faster than neldermead

# scipy\_neldermead, nlopt\_neldermead

- Popular direct search method
- **scipy** does not support bounds
- **nlopt** requires fewer criterion evaluations in most benchmarks
- Almost never the best choice but sometimes not the worst

# scipy\_ls\_lm, scipy\_ls\_trf

- Derivative based optimizers for nonlinear least squares
- Criterion needs structure:  $F(x) = \sum_i f_i(x)^2$
- In estimagic, criterion returns a dict:

```
def sphere_ls(x):  
    # x are the least squares residuals in the sphere function  
    return {"root_contributions": x, "value": x @ x}
```

- **scipy\_ls\_lm** is better for small problems without bounds
- **scipy\_ls\_trf** is better for problems with many parameters

# nag\_dfols, pounders

- Derivative free trust region method for nonlinear least-squares
- Both beat bobyqa for least-squares problems!
- **nag\_dfols** is usually the fastest
- **nag\_dfols** has advanced options to deal with noise
- **pounders** can do criterion evaluations in parallel

# ipopt

- Probably best open source optimizer for nonlinear constraints

# Practice Session

Play with **algorithm** and **algo\_options** (10 min)

# Bounds and constraints

# Preview of practice session

- Solve optimization problem with
  - parameter bounds
  - fixed parameters
  - linear constraints



# Bounds

- Lower and upper bounds on parameters
- Also called box constraints
- Handled by most algorithms
- Need to hold for start parameters
- Guaranteed to hold during entire optimization
- Specification depends on **params** format

# How to specify bounds for array params

```
>>> def sphere(x):  
...     return x @ x  
  
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(3) + 1,  
...     lower_bounds=np.ones(3),  
...     algorithm="scipy_lbfgsb",  
... )  
>>> res.params  
array([1., 1., 1.]
```

- Specify **lower\_bounds** and **upper\_bounds**
- Use **np.inf** or **-np.inf** to represent no bounds

# How to specify bounds for DataFrame params

```
>>> params = pd.DataFrame({  
...     "value": [1, 2, 3],  
...     "lower_bound": [1, 1, 1],  
...     "upper_bound": [3, 3, 3],  
... },  
...     index=["a", "b", "c"],  
... )  
  
>>> def criterion(p):  
...     return (p["value"] ** 2).sum()  
  
>>> em.minimize(criterion, params, algorithm="scipy_lbfgsb")
```

# How to specify bounds for pytree params

```
params = {"x": np.arange(3), "intercept": 3}

def criterion(p):
    return p["x"] @ p["x"] + p["intercept"]

res = em.minimize(
    criterion,
    params=params,
    algorithm="scipy_lbfgsb",
    lower_bounds={"intercept": -2},
)
```

- Enough to specify the subset of params that actually has bounds
- We try to match your bounds with params
- Raise **InvalidBoundsError** in case of ambiguity

# Constraints

- Constraints are conditions on parameters
- Linear constraints
  - $\min_x f(x) \text{ s.t. } l \leq Ax \leq u$
  - $\min_x f(x) \text{ s.t. } Ax = v$
- Nonlinear constraints:
  - $\min_x f(x) \text{ s.t. } c_1(x) = 0, c_2(x) \geq 0$
- "estimagic-constraints":
  - E.g. find probabilities or covariance matrix, fix parameters, ...

# Example: Find valid probabilities

```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.array([0.1, 0.5, 0.4, 4, 5]),  
...     algorithm="scipy_lbfgsb",  
...     constraints=[  
...         "loc": [0, 1, 2],  
...         "type": "probability"  
...     ]],  
... )  
  
>>> res.params  
array([0.33334, 0.33333, 0.33333, -0., 0.] )
```

- Restrict first 3 parameters to be probabilities
  - Between 0 and 1
  - Sum to 1
- But "scipy\_lbfgsb" is unconstrained?!

# What happened

- Estimagic can implement some types of constraints via reparametrization
- Transforms a constrained problem into an unconstrained one
- Constraints must hold in start params
- Guaranteed to hold during entire optimization

# Which constraints can be handled via reparametrization?

- Fixing parameters (simple but useful)
- Finding valid covariance and correlation matrices
- Finding valid probabilities
- Linear constraints (as long as there are not too many)
  - $\min f(x) \text{ s.t. } A_1 x = 0 \text{ and } A_2 x \leq 0$



# Fixing parameters

```
>>> def criterion(params):
...     offset = np.linspace(1, 0, len(params))
...     x = params - offset
...     return x @ x

unconstrained_optimum = [1, 0.8, 0.6, 0.4, 0.2, 0]

>>> res = em.minimize(
...     criterion=criterion,
...     params=np.array([2.5, 1, 1, 1, 1, -2.5]),
...     algorithm="scipy_lbfgsb",
...     constraints={"loc": [0, 5], "type": "fixed"},
... )
>>> res.params
array([ 2.5,  0.8,  0.6,  0.4,  0.2, -2.5])
```

- **loc** selects location 0 and 5 of the parameters
- **type** states that they are fixed

# Linear constraints

```
>>> res = em.minimize(  
...     criterion=criterion,  
...     params=np.ones(6),  
...     algorithm="scipy_lbfgsb",  
...     constraints={  
...         "loc": [0, 1, 2, 3],  
...         "type": "linear",  
...         "lower_bound": 0.95,  
...         "weights": 0.25,  
...     },  
... )  
>>> res.params  
array([ 1.25,  1.05,  0.85,  0.65,  0.2 , -0.])
```

- Impose that average of first 4 parameters is larger than 0.95
- Weights can be scalars or same length as selected parameters
- Use "value" instead of "lower\_bound" for linear equality constraint

# Constraints have to hold

```
>>> em.minimize(  
...     criterion=sphere,  
...     params=np.array([1, 2, 3, 4, 5]),  
...     algorithm="scipy_lbfgsb",  
...     constraints={"loc": [0, 1, 2], "type": "probability"},  
... )
```

InvalidParamsError: A constraint of type 'probability' is not fulfilled in params, please make sure that it holds for the starting values. The problem arose because: Probabilities do not sum to 1. The names of the involved parameters are: ['0', '1', '2'] The relevant constraint is:  
{'loc': [0, 1, 2], 'type': 'probability', 'index': array([0, 1, 2])}.

# Nonlinear constraints

```
>>> res = em.minimize(  
...     criterion=criterion,  
...     params=np.ones(6),  
...     algorithm="scipy_slsqp",  
...     constraints={  
...         "type": "nonlinear",  
...         "loc": [0, 1, 2, 3, 4],  
...         "func": lambda x: np.prod(x),  
...         "value": 1.0,  
...     },  
... )  
>>> res.params  
array([1.31, 1.16, 1.01, 0.87, 0.75, -0.])
```

- Restrict the product of first five params to 1
- Only work with some optimizers
- **func** can be an arbitrary python function of params that returns a number, array or pytree
- Use "lower\_bound" and "upper\_bound" instead of "value" for inequality constraints

# Parameter selection methods

- **"loc"** can be replaced other things
- If params is a DataFrame with "value" column
  - **"query"**: An arbitrary query string that selects the relevant rows
  - **"loc"**: Will be passed to **DataFrame.loc**
- Always
  - **"selector"**: A function that takes params as argument and returns a subset of params
- More in the documentation

# Practice Session

Constrained optimization (15 min)


# Global optimization

# Global vs local optimization


- Local: Find any local optimum
  - All we have done so far
- Global: Find best local optimum
  - Needs bounds to be well defined
  - Extremely challenging in high dimensions
- Local = global for convex problems



# Examples






Alpine



Ackley

# How about brute force?



Dimension	Runtime (1 ms per evaluation, 100 points per dimension)
-----------	---------------------------------------------------------

1	100 ms
---	--------

2	10 s
---	------

3	16 min
---	--------

4	27 hours
---	----------

5	16 weeks
---	----------

6	30 years
---	----------

# Genetic algorithms

- Heuristic inspired by natural selection
- Random initial population of parameters
- In each evolution step:
  - Evaluate "fitness" of all population members
  - Replace worst by combinations of good ones
- Converge when max iterations are reached
- Examples: "**pygmo\_gaco**", "**pygmo\_bee\_colony**", "**nlopt\_crs2\_lm**", ...

# Bayesian optimization

- Evaluate criterion on grid or sample of parameters
- Build surrogate model of criterion
- In each iteration
  - Do new criterion evaluations at promising points
  - Improve surrogate model
- Converge when max iterations is reached

# Multistart optimization:

- Evaluate criterion on random exploration sample
- Run local optimization from best point
- In each iteration:
  - Combine best parameter and next best exploration point
  - Run local optimization from there
- Converge if current best optimum is rediscovered several times

# Multistart example

```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="scipy_neldermead",  
...     soft_lower_bounds=np.full(5, -5),  
...     soft_upper_bounds=np.full(5, 15),  
...     multistart=True,  
...     multistart_options={  
...         "convergence.max_discoveries": 5,  
...         "n_samples": 1000  
...     },  
... )  
>>> res.params  
array([0., 0., 0., 0., 0.]
```

- Turn local optimizers global
- Inspired by tiktok algorithm
- Use any optimizer
- Distinguish hard and soft bounds

# How to choose

- Extremely expensive criterion (i.e. can only do a few evaluations):
  - Bayesian optimization
- Well behaved function:
  - Multistart with local optimizer tailored to function properties
- Rugged function with extremely many local optima
  - Genetic optimizer
  - Consider refining the result with local optimizer
- All are equally parallelizable

# Summary



# Summary

- You have solved a simple problem with `scipy.optimize`
- For larger problems, `estimagic` provides more convenience
  - ``params`` don't have to be flat arrays
  - Support for many more algorithms
  - ``criterion_plot`` and ``params_plot`` help you select the right algorithm
  - Logging and error handling help to deal with crashes
  - Multistart to make local optimizers global
- Many more features to explore in the documentation

# How to contribute

- Make issues or provide feedback
- Improve or extend the documentation
- Suggest, wrap or implement new optimizers
- Teach estimagic to colleagues, students and friends
- Make us happy by giving us a star on [github.com/OpenSourceEconomics/estimagic](https://github.com/OpenSourceEconomics/estimagic)